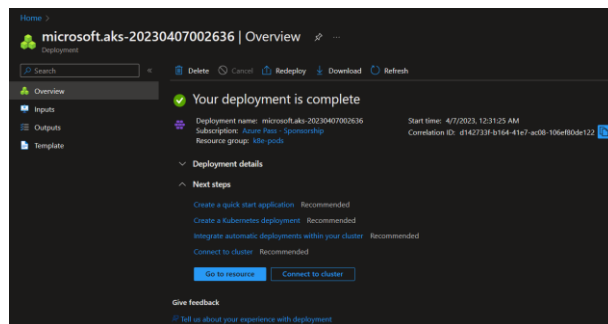


Kubernetes Pods

Practice1: Simple pods operations

1. Login to Azure and connect to your AKS cluster.

- We first need to create AKS cluster. We do that by going to Kubernetes services in Azure and go to + Create and Create Kubernetes Cluster. And we input the requested fields. After that we click on Review + create. We wait for successful validation and then we click create. If the deployment of the AKS cluster is successful, we get the following:



After the cluster is created, we need to connect to that cluster by using Azure CLI. We use the following commands:

```
PS /home/ana> az account set --subscription 9947c7d8-0002-4f95-9be2-e61cef4a15b3
PS /home/ana> az aks get-credentials --resource-group k8e-pods --name kubernetes-pods
Merged "kubernetes-pods" as current context in /home/ana/.kube/config
PS /home/ana>
```

Once we have run the commands above to connect to the cluster, we can run any kubectl commands.

2. Check how many pods run under the default namespace. Run kubectl get pods.

```
PS /home/ana> kubectl get pods
No resources found in default namespace.
PS /home/ana>
```

After we run the command, we can see that there are no pods.

- ### 3. You should not see any pod under the default namespace. Now check all namespaces. Run kubectl get pods --all-namespaces.
- When we run the command, we can see that there are several pods automatically created.

```
PS /home/ana> kubectl get pods --all-namespaces
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  ama-logs-7p62j                         2/2     Running   0           55m
kube-system  ama-logs-rs-6755dd59ff-722qn          1/1     Running   0           19m
kube-system  azure-ip-masq-agent-b7d7d            1/1     Running   0           55m
kube-system  cloud-node-manager-r88mf             1/1     Running   0           55m
kube-system  coredns-59b6bf8b4f-cr99f            1/1     Running   0           19m
kube-system  coredns-59b6bf8b4f-m2qkv            1/1     Running   0           19m
kube-system  coredns-autoscaler-64b6477b8b-8gt6k  1/1     Running   0           19m
kube-system  csi-azuredisk-node-5vqnq             3/3     Running   0           55m
kube-system  csi-azurefile-node-j9h7w            3/3     Running   0           55m
kube-system  connectivity-agent-94874848f-hbnnm4  1/1     Running   0           34m
kube-system  connectivity-agent-94874848f-r6f28    1/1     Running   0           19m
kube-system  kube-proxy-jpmpb                     1/1     Running   0           55m
kube-system  metrics-server-7dd74d8758-bc9vg      2/2     Running   0           18m
kube-system  metrics-server-7dd74d8758-rfw6d      2/2     Running   0           18m
PS /home/ana>
```

4. How many pods do you see? Who deployed these pods? Why are they deployed?
 - We see 14 pods here, but they can vary. The number will depend on the specific Kubernetes cluster and what applications and components are running in it. The kube-system namespace in Kubernetes is reserved for system-level components that are critical to the operation of the Kubernetes cluster. The reason for deploying kube-system pods is to provide essential cluster functionality, including Control Plane Components, Networking, Logging and Monitoring, Security.
5. Now deploy your first pod using the imperative approach. Run `kubectl run nginx --image=nginx`.

```
PS /home/ana> kubectl run nginx --image=nginx
pod/nginx created
PS /home/ana>
```

6. Validate if the pods have been created. What is the status of your pod?

```
PS /home/ana> kubectl run nginx --image=nginx
pod/nginx created
PS /home/ana> kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           81s
PS /home/ana>
```

We validate them with the command `kubectl run nginx --image=nginx`. And we can see that the status of the pod created is **running**.

7. Check the logs coming out of your pod. Run `kubectl logs nginx`.

```
PS /home/ana> kubectl logs nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
2023/04/07 00:17:26 [notice] 1#1: using the "epoll" event method
2023/04/07 00:17:26 [notice] 1#1: nginx/1.23.4
2023/04/07 00:17:26 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2023/04/07 00:17:26 [notice] 1#1: OS: Linux 5.4.0-1104-azure
2023/04/07 00:17:26 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2023/04/07 00:17:26 [notice] 1#1: start worker processes
2023/04/07 00:17:26 [notice] 1#1: start worker process 29
2023/04/07 00:17:26 [notice] 1#1: start worker process 30
/docker-entrypoint.sh: Configuration complete; ready for start up
```

8. Run following command to check current resource consumption of your pod: `kubectl top pod nginx`.

```
PS /home/ana> kubectl top pod nginx
NAME          CPU(cores)   MEMORY(bytes)
nginx         0m           3Mi
PS /home/ana> █
```

9. Check on which Node your pods have been scheduled. Run `kubectl get pods -o wide`.

```
PS /home/ana> kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP          NODE                                     NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0          20m   10.244.1.14 aks-agentpool-34438559-vmss000002    <none>           <none>
PS /home/ana> █
```

10. Try to find the same information but this time running `kubectl describe pod nginx`.

```
PS /home/ana> kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP          NODE                                     NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0          20m   10.244.1.14 aks-agentpool-34438559-vmss000002    <none>           <none>
PS /home/ana> kubectl describe pod nginx
Name:         nginx
Namespace:    default
Priority:      0
Service Account: default
Node:         aks-agentpool-34438559-vmss000002/10.224.0.6
Start Time:   Fri, 07 Apr 2023 00:17:23 +0000
Labels:       run=nginx
Annotations:   <none>
Status:       Running
IP:           10.244.1.14
IPs:          IP: 10.244.1.14
Containers:
  nginx:
    Container ID:  containerd://e9e53f793306d6af681585f14f75b8047fcb68c79a1fcd69105ab212409e9f
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:2ab30d6ac53580a6db8b657abf0f68d75360ff5cc1670a85acb5bd85ba1b19c0
    Port:         <none>
    Host Port:     <none>
    State:         Running
      Started:     Fri, 07 Apr 2023 00:17:26 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-qwbwb (ro)
Conditions:
  Type             Status
  Initialized       True
  Ready            True
  ContainersReady   True
```

11. Delete your pod using `kubectl delete pod nginx`.

```
PS /home/ana> kubectl delete pod nginx
pod "nginx" deleted
PS /home/ana> █
```

12. Let's find the image used on one of the coredns pods under the kube-system namespace.
 13. Once again list all pods under all namespaces.

```
PS /home/ana> kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	ama-logs-rs-6755dd59ff-5rbc2	1/1	Running	0	71m
kube-system	ama-logs-wlnw9	2/2	Running	0	72m
kube-system	azure-ip-masq-agent-xtmns	1/1	Running	0	72m
kube-system	cloud-node-manager-nr54r	1/1	Running	0	72m
kube-system	coredns-59b6bf8b4f-4mfkc	1/1	Running	0	71m
kube-system	coredns-59b6bf8b4f-55jz5	1/1	Running	0	71m
kube-system	coredns-autoscaler-64b6477b8b-mmm9z	1/1	Running	0	71m
kube-system	csi-azuredisk-node-2f456	3/3	Running	0	72m
kube-system	csi-azurefile-node-2ts9t	3/3	Running	0	72m
kube-system	konnectivity-agent-94874848f-8cb2j	1/1	Running	0	71m
kube-system	konnectivity-agent-94874848f-mbjls	1/1	Running	0	71m
kube-system	kube-proxy-7svm9	1/1	Running	0	72m
kube-system	metrics-server-7dd74d8758-2rtx7	2/2	Running	0	71m
kube-system	metrics-server-7dd74d8758-lc98h	2/2	Running	0	71m

```
PS /home/ana> 
```

14. Note one of the coredns pods. Now run kubectl describe pod <coredns-name> -n kube-system. Replace the <coredns-name> place holder with noted name. We will use this coredns pod: **coredns-59b6bf8b4f-4mfkc**

```
PS /home/ana> kubectl describe pod coredns-59b6bf8b4f-4mfkc -n kube-system
```

```
Name: coredns-59b6bf8b4f-4mfkc
Namespace: kube-system
Priority: 2000001000
Priority Class Name: system-node-critical
Service Account: coredns
Node: aks-agentpool-34438559-vmss000003/10.244.0.4
Start Time: Sat, 08 Apr 2023 22:29:09 +0000
Labels: k8s-app=kube-dns
        kubernetes.io/cluster-service=true
        pod-template-hash=59b6bf8b4f
        version=v20
Annotations: prometheus.io/port: 9153
Status: Running
IP: 10.244.0.6
IPs:
  IP: 10.244.0.6
Controlled By: ReplicaSet/coredns-59b6bf8b4f
Containers:
  coredns:
    Container ID: containerd://6737f21aaedb31bb3521389e6a1cb7c458d2a480376752fc2e817c06651cffa4
    Image: mcr.microsoft.com/oss/kubernetes/coredns:v1.9.3
    Image ID: sha256:c38f956b642366c8eeb0babfda6b0bb2aa92f27a968589804cadb445f6df72d6
    Ports: 53/UDP, 53/TCP, 9153/TCP
    Host Ports: 0/UDP, 0/TCP, 0/TCP
    Args:
      -conf
      /etc/coredns/Corefile
    State: Running
      Started: Sat, 08 Apr 2023 22:29:13 +0000
    Ready: True
    Restart Count: 0
    Limits:
```

```

Limits:
  cpu:      3
  memory:   500Mi
Requests:
  cpu:      100m
  memory:   70Mi
Liveness:   http-get http://:8080/health delay=60s timeout=5s period=10s #success=1 #failure=5
Readiness:  http-get http://:8181/ready delay=0s timeout=1s period=10s #success=1 #failure=3
Environment:
  KUBERNETES_PORT_443_TCP_ADDR:  kubernetes-pods-dns-djj2t7xv.hcp.eastus.azmk8s.io
  KUBERNETES_PORT:                tcp://kubernetes-pods-dns-djj2t7xv.hcp.eastus.azmk8s.io:443
  KUBERNETES_PORT_443_TCP:        tcp://kubernetes-pods-dns-djj2t7xv.hcp.eastus.azmk8s.io:443
  KUBERNETES_SERVICE_HOST:        kubernetes-pods-dns-djj2t7xv.hcp.eastus.azmk8s.io
Mounts:
  /etc/coredns from config-volume (ro)
  /etc/coredns/custom from custom-config-volume (ro)
  /tmp from tmp (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-kh54z (ro)
Conditions:
  Type              Status
  Initialized        True
  Ready              True
  ContainersReady    True
  PodScheduled       True
Volumes:
  config-volume:
    Type:          ConfigMap (a volume populated by a ConfigMap)
    Name:           coredns
    Optional:       false
  custom-config-volume:
    Type:          ConfigMap (a volume populated by a ConfigMap)
    Name:           coredns-custom
    Optional:       true

```

```

tmp:
  Type:          EmptyDir (a temporary directory that shares a pod's lifetime)
  Medium:
  SizeLimit:     <unset>
kube-api-access-kh54z:
  Type:          Projected (a volume that contains injected data from multiple sources)
  TokenExpirationSeconds:  3607
  ConfigMapName:          kube-root-ca.crt
  ConfigMapOptional:      <nil>
  DownwardAPI:            true
QoS Class:           Burstable
Node-Selectors:      <none>
Tolerations:         CriticalAddonsOnly op=Exists
                     node.kubernetes.io/master:NoSchedule
                     node.kubernetes.io/memory-pressure:NoSchedule op=Exists
                     node.kubernetes.io/not-ready:NoExecute op=Exists for 30s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 30s
Events:
PS /home/ana>

```

15. Inspect the output and locate the image information.

```

Containers:
  coredns:
    Container ID:  containerd://6737f21aaedb31bb3521389e6a1cb7c458d2a480376752fc2e817c06651cffa4
    Image:         mcr.microsoft.com/oss/kubernetes/coredns:v1.9.3
    Image ID:      sha256:c38f956b642366c8eeb0babfda6b0bb2aa92f27a968589804cadb445f6df72d6
    Ports:         53/UDP, 53/TCP, 9153/TCP
    Host Ports:    0/UDP, 0/TCP, 0/TCP
    Args:
      -conf
      /etc/coredns/Corefile
    State:         Running
      Started:     Sat, 08 Apr 2023 22:29:13 +0000
    Ready:         True
    Restart Count: 0

```

- Now let us check the logs of the metrics-server pod. Run the same command as in step 7 but don't forget to add the namespace in which this pod is created.

We use the command **kubectl logs metrics-server-7dd74d8758-lc98h --container metrics-server --namespace kube-system** where we need to specify the container name as well as the namespace, otherwise if we don't specify the namespace and the container we would receive an error.

```
PS /home/ana> kubectl logs metrics-server-7dd74d8758-lc98h --container metrics-server --namespace kube-system
I0408 22:29:38.839288      1 serving.go:342] Generated self-signed cert (/tmp/apiserver.crt, /tmp/apiserver.key)
I0408 22:29:47.736235      1 secure_serving.go:266] Serving securely on [::]:4443
I0408 22:29:47.736662      1 requestheader_controller.go:169] Starting RequestHeaderAuthRequestController
I0408 22:29:47.736677      1 shared_informer.go:240] Waiting for caches to sync for RequestHeaderAuthRequestController
W0408 22:29:47.736700      1 shared_informer.go:372] The sharedIndexInformer has started, run more than once is not allowed
I0408 22:29:47.737350      1 dynamic_serving_content.go:131] "Starting controller" name="serving-cert::/tmp/apiserver.crt:/tmp/apiserver.key"
I0408 22:29:47.737573      1 configmap_cafile_content.go:201] "Starting controller" name="client-ca::kube-system:extension-apiserver-authentication::client-ca-file"
I0408 22:29:47.737582      1 shared_informer.go:240] Waiting for caches to sync for client-ca::kube-system:extension-apiserver-authentication::client-ca-file
I0408 22:29:47.737598      1 configmap_cafile_content.go:201] "Starting controller" name="client-ca::kube-system:extension-apiserver-authentication::requestheader-client-ca-file"
I0408 22:29:47.737602      1 shared_informer.go:240] Waiting for caches to sync for client-ca::kube-system:extension-apiserver-authentication::requestheader-client-ca-file
I0408 22:29:47.738459      1 tlsconfig.go:240] "Starting DynamicServingCertificateController"
I0408 22:29:47.840654      1 shared_informer.go:247] Caches are synced for client-ca::kube-system:extension-apiserver-authentication::requestheader-client-ca-file
I0408 22:29:47.840699      1 shared_informer.go:247] Caches are synced for RequestHeaderAuthRequestController
I0408 22:29:47.840757      1 shared_informer.go:247] Caches are synced for client-ca::kube-system:extension-apiserver-authentication::client-ca-file
PS /home/ana>
```

Errors when either of the container or the namespace are not specified:

```
PS /home/ana> kubectl logs metrics-server-7dd74d8758-lc98h --container metrics-server
Error from server (NotFound): pods "metrics-server-7dd74d8758-lc98h" not found
PS /home/ana> kubectl logs metrics-server-7dd74d8758-lc98h
Error from server (NotFound): pods "metrics-server-7dd74d8758-lc98h" not found
PS /home/ana> kubectl logs metrics-server-7dd74d8758-lc98h -n kube-system
Defaulted container "metrics-server-vpa" out of: metrics-server-vpa, metrics-server
ERROR: logging before flag.Parse: I0408 22:29:15.857596      1 pod_nanny.go:68] Invoked by [/pod_nanny --config-dir=/etc/config --cpu=44m --extra-cpu=0.5m --memory=51Mi --extra-memory=4Mi --poll-period=180000 --threshold=5 --deployment=metrics-server --container=metrics-server]
ERROR: logging before flag.Parse: I0408 22:29:15.857816      1 pod_nanny.go:69] Version: 1.8.14
ERROR: logging before flag.Parse: I0408 22:29:15.857950      1 pod_nanny.go:85] Watching namespace: kube-system, pod: metrics-server-7dd74d8758-lc98h, container: metrics-server.
ERROR: logging before flag.Parse: I0408 22:29:15.858070      1 pod_nanny.go:86] storage: MISSING, extra_storage: 0Gi
ERROR: logging before flag.Parse: I0408 22:29:16.548364      1 pod_nanny.go:189] Failed to read data from config file "/etc/config/NannyConfiguration": open /etc/config/NannyConfiguration: no such file or directory, using default parameters
ERROR: logging before flag.Parse: I0408 22:29:16.548620      1 pod_nanny.go:116] cpu: 44m, extra_cpu: 0.5m, memory: 51Mi, extra_memory: 4Mi
ERROR: logging before flag.Parse: I0408 22:29:16.548643      1 pod_nanny.go:145] Resources: [{Base:{i:{value:44 scale:-3} d:{Dec:nil} s:44m Format:DecimalSI} ExtraPerNode:{i:{value:5 scale:-4} d:{Dec:nil} s: Format:DecimalSI} Name:cpu} {Base:{i:{value:53477376 scale:0} d:{Dec:nil} s:51Mi Format:BinarySI} ExtraPerNode:{i:{value:4194304 scale:0} d:{Dec:nil} s:4Mi Format:BinarySI} Name:memory}]
ERROR: logging before flag.Parse: E0408 22:29:26.570191      1 nanny_lib.go:128] Get "https://kubernetes-pods-dns-djj2t7xv.hcp.eastus.azurek8s.io:443/api/v1/nodes?resourceVersion=0": dial tcp: lookup kubernetes-pods-dns-djj2t7xv.hcp.eastus.azurek8s.io on 10.0.0.10:53: read udp 10.244.0.5:37713->10.0.0.10:53: read: connection refused
PS /home/ana>
```

Practice2: Working with pod manifest files

- Now it is time to deploy pod using manifest file (declarative approach). Copy the following code block on your local computer in a file called redis.yaml:
- Try to deploy the pod defined in redis.yaml. Run **kubectl create -f redis.yaml**.

```
PS /home/ana> kubectl create -f redis.yaml
error: resource mapping not found for name: "" namespace: "" from "redis.yaml": no matches for kind "pod" in version "v1"
ensure CRDs are installed first
PS /home/ana>
```


3. You will receive errors on your screen. Your next task will be to correct the syntax of the code you just copied. You can use the online Kubernetes documentation, or you can search the internet in general.

- This is the corrected syntax of the redis.yaml file.

In yaml files everything is indented. In the given code there wasn't any indentation. Next the apiVersion was wrong, it should have been v1, because v11 doesn't exist. In kind, pod needed to be capitalized.

```
redis.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: static-web
5    labels:
6      role: myrole
7  spec:
8    containers:
9      - name: redis
10       image: redis123
11
```

4. When you solve all the syntax errors your pod should be deployed but is it running? What is the status of your pod?

- The status of the pod here is ErrImagePull. This means that kubernetes is unable to locate the image.

```
PS /home/ana> kubectl create -f redis.yaml
pod/static-web created
PS /home/ana> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    0/1     ErrImagePull  0          107s
PS /home/ana>
```

5. Check the events associated with this pod. Run the **kubectl describe pod static-web** command. What are the events showing? Why is your pod not running?

(This is just the events part)

```
Events:
  Type     Reason      Age    From          Message
  ----     -
  Normal   Scheduled   18m    default-scheduler   Successfully assigned default/static-web to aks-agentpool-34438559-vmss000004
  Normal   Pulling     17m    kubelet        Pulling image "redis123"
  Warning   Failed      17m    kubelet        Failed to pull image "redis123": rpc error: code = Unknown desc = failed to pull and unpack image "docker.io/library/redis123:latest": failed to resolve reference "docker.io/library/redis123:latest": pull access denied, repository does not exist or may require authorization: server message: insufficient_scope: authorization failed
  Warning   Failed      17m    kubelet        Error: ErrImagePull
  Warning   Failed      16m    kubelet        Error: ImagePullBackOff
  Normal   BackOff     3m35s  kubelet        Back-off pulling image "redis123"
PS /home/ana>
```

- The status of the pod here is ErrImagePull. This means that kubernetes is unable to locate the image.
6. Find the correct image (check the Docker hub page) and correct it in the manifest.
- The correct image was redis:latest instead redis123.

```
redis.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: static-web
5    labels:
6      role: myrole
7  spec:
8    containers:
9      - name: redis
10       image: redis:latest
```

7. Locate the image information and put the correct image name. Redeploy the pod (first run **kubect1 delete pod static-web** to delete the pod, then run kubect1 create once again).
8. Check the status of your pod. It should be running now.

```
PS /home/ana> kubect1 get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    1/1     Running   0           59s
PS /home/ana>
```

9. Now you can delete the pod. Try to delete it using the **kubect1 delete -f redis.yaml**.

```
PS /home/ana> kubect1 delete -f redis.yaml
pod "static-web" deleted
PS /home/ana>
```

10. Your next task is to create and test nginx pod definition. Your definition should use the nginx official image, should use label named app with value frontend and should publish port 80. Make sure you complete this task because we will use this template in our next Labs. Your nginx pod should be running without any issues.

```
nginx.yml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx-pod
5    labels:
6      app: frontend
7  spec:
8    containers:
9      - name: nginx-container
10        image: nginx:latest
11        ports:
12          - containerPort: 80
13
```

```
PS /home/ana> kubect1 create -f nginx.yml
pod/nginx-pod created
PS /home/ana> kubect1 get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx-pod     1/1     Running   0           6s
PS /home/ana>
```

11. Final task of this practice will be to define pod definition with following details:

```
nginx-pod.yml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: memcached-pod
5    labels:
6      app: web
7  spec:
8    restartPolicy: Never
9    containers:
10      - name: memcached-container
11        image: memcached
12        ports:
13          - containerPort: 11211
14        resources:
15          requests:
16            cpu: 0.35
17            memory: 150Mi
18          limits:
19            cpu: 0.5
20            memory: 250Mi
21
```

- Image=memcached
- Port= 11211
- Label app=web
- CPU request=0.35 cores
- RAM request=0.15 GB
- CPU limit=0.5 cores
- Ram limit=0.25 GB
- Restart policy=Never

12. Don't forget to try your pod definition.

```
PS /home/ana> kubect1 create -f nginx-pod.yml
pod/memcached-pod created
PS /home/ana> kubect1 get pods
NAME          READY   STATUS    RESTARTS   AGE
memcached-pod 1/1     Running   0           5s
PS /home/ana>
```


Practice3: Multi-container pods

1. Once finished you can try to create multi-container pod definition. Your multi-container pod should use redis and nginx containers with port 6379 and 80 published respectively. Label name should be app with value web.
2. Note that in reality there is no sense to put the redis and nginx under the same pod but it can be done for the purpose of learning.
3. Deploy your multi-container pod. It should have running status.

What is written under Ready column when you **kubectl get pods**? – 0/2 in the first running of **kubectl get pods**, and 2/2 when we run **kubectl get pods** for the second time.

Why your pod displays different values for ready? – When the command is ran for the first time the containers were still starting up, but for the second time the command is ran the containers had finished starting up.

```
PS /home/ana> kubectl create -f redis-nginx-pod.yml
pod/redis-nginx-pod created
PS /home/ana> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
memcached-pod	1/1	Running	0	19m
redis-nginx-pod	0/2	Pending	0	9s

```
PS /home/ana> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
redis-nginx-pod	2/2	Running	0	42s

4. Kubectl describe your new pod, and locate the containers section. How many containers are listed?
 - To describe our new pod we use **kubectl describe pod redis-nginx-pod**. And there are 2 containers.

```
Containers:
  redis-container:
    Container ID:  containerd://0bbe35a6fd7e0e886877f25b289d73115228ccd66fb612e1d3fd8a0023da502c
    Image:          redis
    Image ID:       docker.io/library/redis@sha256:7b83a0167532d4320a87246a815a134e19e31504d85e8e55f0bb5bb9edf70448
    Port:          6379/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Sun, 09 Apr 2023 22:12:28 +0000
    Ready:         True
    Restart Count: 0
    Limits:
      cpu:         500m
      memory:      250Mi
    Requests:
      cpu:         350m
      memory:      150Mi
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-8ngnm (ro)
  nginx-container:
    Container ID:  containerd://a7ae6fcaa15a8d552751c2b57aa6e6de535fb2b966afed53553bf94e367cc7ad
    Image:          nginx
    Image ID:       docker.io/library/nginx@sha256:2ab30d6ac53580a6db8b657abf0f68d75360ff5cc1670a85acb5bd85ba1b19c0
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Sun, 09 Apr 2023 22:12:29 +0000
    Ready:         True
    Restart Count: 0
    Limits:
      cpu:         500m
      memory:      250Mi
    Requests:
```

5. Delete all the pods under the default namespace.
6. Don't delete any of the manifest files you have created so far.
 - Before deleting the pod we first list all the pods in the default namespace, and confirm that the redis-nginx-pod pod is present. After we confirm that the redis-nginx-pod pod is present we delete it.
 - **kubectl delete -f redis-nginx-pod.yml** command does not delete the manifest file we created.

```

PS /home/ana> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
redis-nginx-pod 2/2     Running   0           31m
PS /home/ana> kubectl delete -f redis-nginx-pod.yml
pod "redis-nginx-pod" deleted
PS /home/ana> kubectl get pods
No resources found in default namespace.
PS /home/ana> ls
cloudrive  Microsoft  nginx-pod.yml  nginx.yml  redis-nginx-pod.yml  redis.yaml
PS /home/ana>

```

Practice4: Probes

3. Run `kubectl create -f probes_exec.yaml`.

```

PS /home/ana> kubectl create -f probes_exec.yaml
pod/liveness-exec created
PS /home/ana>

```

4. Run `kubectl describe pod liveness-exec` immediately after you deploy the pod. The output should indicate that no liveness probes have failed yet.

```

Events:
  Type    Reason      Age   From              Message
  ----    -
Normal   Scheduled   3s    default-scheduler Successfully assigned default/liveness-exec to aks-agentpool-34438559-vmss000005
Normal   Pulling     3s    kubelet           Pulling image "k8s.gcr.io/busybox"
Normal   Pulled      3s    kubelet           Successfully pulled image "k8s.gcr.io/busybox" in 288.282835ms
Normal   Created     3s    kubelet           Created container liveness
Normal   Started     2s    kubelet           Started container liveness

```

5. After 35 seconds, view the Pod events again. Run `kubectl describe pod liveness-exec`.

6. At the bottom of the output, there should be a messages indicating that the liveness probes have failed, and the containers have been killed and recreated.

```

Events:
  Type    Reason      Age          From              Message
  ----    -
Normal   Scheduled   54s          default-scheduler Successfully assigned default/liveness-exec to aks-agentpool-34438559-vmss000005
Normal   Pulling     54s          kubelet           Pulling image "k8s.gcr.io/busybox"
Normal   Pulled      54s          kubelet           Successfully pulled image "k8s.gcr.io/busybox" in 288.282835ms
Normal   Created     54s          kubelet           Created container liveness
Normal   Started     53s          kubelet           Started container liveness
Warning  Unhealthy   9s (x3 over 19s) kubelet           Liveness probe failed: cat: can't open '/tmp/healthcheck': No such file or directory
Normal   Killing     9s           kubelet           Container liveness failed liveness probe, will be restarted

```

7. Wait another 30 seconds, and verify that the container has been restarted. Run `kubectl get pod liveness-exec`.

```

Events:
  Type      Reason      Age           From          Message
  ----      -
  Normal    Scheduled   91s           default-scheduler Successfully assigned default/liveness-exec to aks-agentpool-34438559-vmss000005
  Normal    Pulled      91s           kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 288.282835ms
  Warning   Unhealthy   46s (x3 over 56s) kubelet        Liveness probe failed: cat: can't open '/tmp/healthz': No such file or directory
  Normal    Killing     46s           kubelet        Container liveness failed liveness probe, will be restarted
  Normal    Pulling     16s (x2 over 91s) kubelet        Pulling image "k8s.gcr.io/busybox"
  Normal    Pulled      16s           kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 298.73071ms
  Normal    Created     15s (x2 over 91s) kubelet        Created container liveness
  Normal    Started     15s (x2 over 90s) kubelet        Started container liveness
PS /home/ana> 

```

```

Events:
  Type      Reason      Age           From          Message
  ----      -
  Normal    Scheduled   4m6s          default-scheduler Successfully assigned default/liveness-exec to aks-agentpool-34438559-vmss000005
  Normal    Pulled      4m6s          kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 288.282835ms
  Normal    Pulled      2m51s         kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 298.73071ms
  Normal    Created     95s (x3 over 4m6s) kubelet        Created container liveness
  Normal    Started     95s (x3 over 4m5s) kubelet        Started container liveness
  Normal    Pulled      95s           kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 343.468964ms
  Warning   Unhealthy   51s (x9 over 3m31s) kubelet        Liveness probe failed: cat: can't open '/tmp/healthz': No such file or directory
  Normal    Killing     51s (x3 over 3m21s) kubelet        Container liveness failed liveness probe, will be restarted
  Normal    Pulling     21s (x4 over 4m6s) kubelet        Pulling image "k8s.gcr.io/busybox"
PS /home/ana> 

```

8. The output should show that RESTARTS has been incremented.

```

PS /home/ana> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
liveness-exec 1/1     Running   1 (49s ago) 2m4s
PS /home/ana> 

```

11. For the first 10 seconds that the container is alive, the /healthz handler returns a status of 200. After that, the handler returns a status of 500.

```

Events:
  Type      Reason      Age   From          Message
  ----      -
  Normal    Scheduled   6s     default-scheduler Successfully assigned default/liveness-http to aks-agentpool-34438559-vmss000005
  Normal    Pulling     5s     kubelet        Pulling image "k8s.gcr.io/liveness"
  Normal    Pulled      5s     kubelet        Successfully pulled image "k8s.gcr.io/liveness" in 226.775943ms
  Normal    Created     5s     kubelet        Created container liveness
  Normal    Started     5s     kubelet        Started container liveness

```

13. Immediately run (you only have 10 secs to run this command) `kubectl describe pod liveness-http`.

14. Your pod should be live and running.

Events:				
Type	Reason	Age	From	Message
-----	-----	----	----	-----
Normal	Scheduled	14s	default-scheduler	Successfully assigned default/liveness-http to aks-agentpool-34438559-vmss000005
Normal	Pulling	13s	kubelet	Pulling image "k8s.gcr.io/liveness"
Normal	Pulled	13s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 226.775943ms
Normal	Created	13s	kubelet	Created container liveness
Normal	Started	13s	kubelet	Started container liveness
Warning	Unhealthy	1s	kubelet	Liveness probe failed: HTTP probe failed with statuscode: 500

15. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again **kubectl describe pod liveness-http**.

Events:				
Type	Reason	Age	From	Message
-----	-----	----	----	-----
Normal	Scheduled	26s	default-scheduler	Successfully assigned default/liveness-http to aks-agentpool-34438559-vmss000005
Normal	Pulled	25s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 226.775943ms
Normal	Pulling	7s (x2 over 25s)	kubelet	Pulling image "k8s.gcr.io/liveness"
Normal	Created	7s (x2 over 25s)	kubelet	Created container liveness
Normal	Started	7s (x2 over 25s)	kubelet	Started container liveness
Warning	Unhealthy	7s (x3 over 13s)	kubelet	Liveness probe failed: HTTP probe failed with statuscode: 500
Normal	Killing	7s	kubelet	Container liveness failed liveness probe, will be restarted
Normal	Pulled	7s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 253.777833ms

16. You should see the same output as in step 7. Kubelet will reboot the container.

17.

18. Run **kubectl create -f probes_tcp.yaml**.

Events:				
Type	Reason	Age	From	Message
-----	-----	----	----	-----
Normal	Scheduled	5s	default-scheduler	Successfully assigned default/liveness-tcp to aks-agentpool-34438559-vmss000005
Normal	Pulled	5s	kubelet	Container image "k8s.gcr.io/goproxy:0.1" already present on machine
Normal	Created	4s	kubelet	Created container goproxy
Normal	Started	4s	kubelet	Started container goproxy

19. Immediately run (you only have 10 secs to run this command) **kubectl describe pod liveness-tcp**.

20.

21. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again **kubectl describe pod liveness-tcp**.

22.

23.

24.

25. Run **kubectl create -f readiness_http.yaml**.

```
PS /home/ana> kubectl create -f readiness_http.yaml
pod/readiness-http created
PS /home/ana> 
```

26. Run **kubectl get pods -A** to see the status of your pod.

27. Pods and their status and ready states will be displayed; our pod should be in running state.

```
PS /home/ana> kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-exec	1/1	Running	20 (5m40s ago)	62m
default	liveness-http	0/1	CrashLoopBackOff	17 (3m8s ago)	39m
default	liveness-tcp	1/1	Running	0	25m
default	readiness-http	1/1	Running	0	2m13s
kube-system	ama-logs-bqztc	2/2	Running	0	153m
kube-system	ama-logs-rs-6755dd59ff-cv14w	1/1	Running	0	135m
kube-system	azure-ip-masq-agent-bmbd9	1/1	Running	0	153m
kube-system	cloud-node-manager-vqvt4	1/1	Running	0	153m
kube-system	coredns-59b6bf8b4f-jd17r	1/1	Running	0	135m
kube-system	coredns-59b6bf8b4f-sgf9q	1/1	Running	0	135m
kube-system	coredns-autoscaler-64b6477b8b-rpjxj	1/1	Running	0	135m
kube-system	csi-azuredisk-node-7wl8t	3/3	Running	0	153m
kube-system	csi-azurefile-node-rlxpw	3/3	Running	0	153m
kube-system	konnnectivity-agent-94874848f-6q9hj	1/1	Running	0	135m
kube-system	konnnectivity-agent-94874848f-ntfgk	1/1	Running	0	135m
kube-system	kube-proxy-64298	1/1	Running	0	153m
kube-system	metrics-server-7dd74d8758-4p4cs	2/2	Running	0	132m
kube-system	metrics-server-7dd74d8758-ftvk4	2/2	Running	0	132m

```
PS /home/ana>
```

28. Run **kubectl describe pod readiness-http**. Examine the events for this pod. Everything should be OK.

```
Events:
```

Type	Reason	Age	From	Message
Normal	Scheduled	6m56s	default-scheduler	Successfully assigned default/readiness-http to aks-agentpool-34438559-vmss000005
Normal	Pulling	6m55s	kubelet	Pulling image "nginx"
Normal	Pulled	6m55s	kubelet	Successfully pulled image "nginx" in 192.41402ms
Normal	Created	6m55s	kubelet	Created container nginx
Normal	Started	6m55s	kubelet	Started container nginx

```
PS /home/ana>
```

29. Now delete the pod and edit the `readiness_http.yaml` so that the port parameter has 81 value.

```
PS /home/ana> kubectl delete -f readiness_http.yaml
pod "readiness-http" deleted
PS /home/ana>
```

30. Run again **kubectl create -f readiness_http.yaml**.

```
PS /home/ana> kubectl create -f readiness_http.yaml
pod/readiness-http created
PS /home/ana>
```

31. Run **kubectl get pods -A** to see the status of your pod. You should see that the pod is running but it is not in ready state.

```
PS /home/ana> kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-exec	0/1	CrashLoopBackOff	23 (119s ago)	73m
default	liveness-http	0/1	CrashLoopBackOff	21 (3m8s ago)	50m
default	liveness-tcp	1/1	Running	0	36m
default	readiness-http	0/1	Running	0	52s
kube-system	ama-logs-bqzfc	2/2	Running	0	164m
kube-system	ama-logs-rs-6755dd59ff-cv14w	1/1	Running	0	146m
kube-system	azure-ip-masq-agent-bmbd9	1/1	Running	0	164m
kube-system	cloud-node-manager-vqvt4	1/1	Running	0	164m
kube-system	coredns-59b6bf8b4f-jdl7r	1/1	Running	0	146m
kube-system	coredns-59b6bf8b4f-sgf9q	1/1	Running	0	146m
kube-system	coredns-autoscaler-64b6477b8b-rpjxj	1/1	Running	0	146m
kube-system	csi-azuredisk-node-7wl8t	3/3	Running	0	164m
kube-system	csi-azurefile-node-rlxpw	3/3	Running	0	164m
kube-system	konnnectivity-agent-94874848f-6q9hj	1/1	Running	0	146m
kube-system	konnnectivity-agent-94874848f-ntfgk	1/1	Running	0	146m
kube-system	kube-proxy-64298	1/1	Running	0	164m
kube-system	metrics-server-7dd74d8758-4p4cs	2/2	Running	0	143m
kube-system	metrics-server-7dd74d8758-ftvk4	2/2	Running	0	143m

```
PS /home/ana>
```

32. Describe the pod. Run **kubectl describe pod readiness-http**.

```
Events:
```

Type	Reason	Age	From	Message
Normal	Scheduled	2m34s	default-scheduler	Successfully assigned default/readiness-http to aks-agentpool-34438559-vmss000005
Normal	Pulling	2m34s	kubelet	Pulling image "nginx"
Normal	Pulled	2m33s	kubelet	Successfully pulled image "nginx" in 169.183015ms
Normal	Created	2m33s	kubelet	Created container nginx
Normal	Started	2m33s	kubelet	Started container nginx
Warning	Unhealthy	114s (x21 over 2m32s)	kubelet	Readiness probe failed: Get "http://10.244.1.22:81/": dial tcp 10.244.1.22:81: connect: connection refused

```
PS /home/ana>
```

33. From the events we can see that readiness probe failed due to the connection being refused therefore pod will not receive any traffic.
34. Delete all pods under the default namespace.
35. Don't delete any manifest files created so far.

```
PS /home/ana> kubectl get pods
No resources found in default namespace.
PS /home/ana> ls
clouddrive  nginx-pod.yml  probes_exec.yml  probes_tcp.yml  redis-nginx-pod.yml
Microsoft  nginx.yml      probes_http.yml  readiness_http.yml  redis.yaml
PS /home/ana>
```